



Task 4 - style-conditioned pix2pix-style cGAN



Generator $G(x, s)$: U-Net, 42.1 M parameters (PatchGAN D : 2.8 M); the style embedding enters the input AND the bottleneck



Spatial augmentation (flip, random crop-resize) uses ONE random draw applied to both photo and sketch